

JDBC

Devdatta Kulkarni

Databases and Java

- JDBC
 - Java Database Connectivity
 - MySQL Connectors
 - <http://dev.mysql.com/downloads/connector/>
- Example
 - JDBC

JDBC topics

- Maven dependency
 - commons-dbcp
 - mysql-connector-java
- Main Concepts/Classes
 - Data source
 - Connection
 - PreparedStatement
 - ResultSet
 - Transactions

Data Source

- A *data source* is where the data is stored by your application (or is accessed from)
 - E.g.: DBMS, a legacy file system
- A JDBC application connects to a target data source using one of the following classes
 - DriverManager
 - Automatically loads any JDBC 4.0 drivers available on the class path
 - DataSource
 - Abstracts the details of the underlying data source from your application

DriverManager

- Database connection is established using the DriverManager's "getConnection" method
 - dbURL =
"jdbc:mysql://localhost:3306/student_courses";
 - conn =
`DriverManager.getConnection(dbURL, "devdatta",
"");`

DataSource

// Setup data source

```
BasicDataSource ds = new BasicDataSource();  
ds.setUsername(this.dbUsername);  
ds.setPassword(this.dbPassword);  
ds.setUrl(this.dbURL);  
ds.setDriverClassName("com.mysql.jdbc.Driver");
```

// Open connection

```
Connection conn = ds.getConnection();
```

Connection class

- Connection
 - Represents a database connection
 - Typical required parameters: Database URL, username, password
- Connection object is *not* thread safe
- How to support application's multiple threads?
 - Option 1: Accessing Connection through synchronized methods
 - Not a good strategy
 - Different threads may be accessing different data items; such threads may end up blocking each other unnecessarily
 - Option 2: Use a separate connection per thread
 - Not a good strategy either
 - Each thread typically represents a user request; your application will have lots of users
 - It cannot open new Connections for all (it will run out of resources)

Connection Pool

- Connection Pooling
 - Keep a pool of open connections
 - Each thread will get its own connection
 - Threads will block if there are no connections available in the pool
 - A connection is given back to the pool once a thread is done using it
 - Apache DBCP
 - Library that implements connection pooling
 - <http://commons.apache.org/proper/commons-dbcp/>

PreparedStatement

- Represents a *parameterized SQL query*
- Parameter values are set based on data provided at runtime
- Use “executeQuery” to execute a statement that returns some value
- Use “executeUpdate” to execute a statement that does not return a value

PreparedStatement

```
String query = "select * from courses where  
course_id=?";
```

```
Connection conn = ds.getConnection();
```

```
PreparedStatement s =  
conn.prepareStatement(query);
```

```
s.setString(1, String.valueOf(courseId));
```

```
ResultSet r = s.executeQuery();
```

PreparedStatement

- “Parameter binding provides a means of separating executable code, such as SQL, from content, transparently handling content encoding and escaping.” [1]

[\[1\] http://martinfowler.com/articles/web-security-basics.html#BindParametersForDatabaseQueries](http://martinfowler.com/articles/web-security-basics.html#BindParametersForDatabaseQueries)

PreparedStatement – Getting keys back

```
String insert = "INSERT INTO courses(name,  
course_num) VALUES(?, ?)";
```

```
PreparedStatement stmt =  
conn.prepareStatement(insert,  
Statement.RETURN_GENERATED_KEYS);
```

```
stmt.setString(1, c.getName());  
stmt.setString(2, c.getCourseNum());
```

```
int affectedRows = stmt.executeUpdate();
```

ResultSet

- Represents a set of records from the table that satisfy the specified query criteria
- The records are accessed through a *cursor*
- The cursor is a pointer that points to one row of data in the ResultSet
- Initially the cursor is positioned before the first row
- The method `ResultSet.next` moves the cursor to the next row

ResultSet – Returning a single row

```
ResultSet r = s.executeQuery();
```

```
    if (!r.next()) {  
        return null;  
    }
```

```
NewCourse c = new NewCourse();  
c.setCourseNum(r.getString("course_num"));  
c.setName(r.getString("name"));  
c.setCourseId(r.getInt("course_id"));
```

ResultSet – Returning multiple rows

```
ResultSet r = s.executeQuery();
```

```
List<NewCourse> courseList = new  
ArrayList<NewCourse>();
```

```
while (!r.next()) {  
    NewCourse c = new NewCourse();  
    c.setCourseNum(r.getString("course_num"));  
    c.setName(r.getString("name"));  
    c.setCourseId(r.getInt("course_id"));  
    courseList.add(c);  
}  
return courseList;
```

ResultSet Types

- Determines how does the *cursor* moves
 - TYPE_FORWARD_ONLY
 - Cursor only moves forward
 - TYPE_SCROLL_INSENSITIVE
 - The result can be scrolled; the result set is insensitive to changes made to the underlying data source while it is open
 - TYPE_SCROLL_SENSITIVE
 - The result can be scrolled; the result set reflects the changes made to the underlying data source while it is open

Transactions

- Problem:
 - How to ensure that more than one database statements are executed as a unit?
- Solution:
 - Transactions

Example where a transaction is needed

- Suppose we want to create a new project and add a meeting as part of project creation:

POST /myeavesdrop/projects/

<project>

 <name>NewSolum</name>

 <description>NewSolum</description>

 <meetings>

 <meeting>

 <name>M1</name>

 <year>2016</year>

 <month>March</month>

 <day>26</day>

 </meeting>

 </meetings>

</project>

Transactions

- We need to define a transaction that includes the following two statements:
 - Insert the new project's name and description into the “projects” table
 - Insert the project's meeting information into the “meetings” table
- We need to execute both these within a single transaction
 - Either both are committed successfully or both are not
 - The execution of the two statements needs to be *'atomic'*

Transactions

- A transaction is a set of statements that are executed as a *unit (also called the 'atomicity' property of a transaction)*
 - Either all statements in a transaction are completed successfully or none are
 - The effect of executing a statement is reflected within a database only after it is *committed*
- A *Connection* is in auto-commit mode by default
 - Each individual statement is treated as a transaction
 - If we want to group two or more statements into a transaction then we have to *disable* the auto-commit mode
 - `conn.setAutoCommit(false);`
 - Commit/rollback of the transaction is our responsibility
 - `conn.commit();`
 - `conn.rollback();`
- <https://docs.oracle.com/javase/tutorial/jdbc/basics/transactions.html>

Transactions

- A simple approach for the database system to guarantee atomicity (*all or nothing property*) is to perform only one transaction at a time
- But this approach would limit performance severely
- So we want transactions to be executed concurrently to gain performance benefits
- However, then we need to solve for *isolation*

Isolation

- What is isolation?
 - At high-level it means preserving data integrity by ensuring that concurrently executing transactions do not interfere with one another
 - What does this mean?
 - Prevent transactions from interfering with one another
 - For example, if two transactions are modifying the same row then the *isolation* property guarantees that none of the transactions would see an *intermediate* state of the row

How to provide isolation?

- We know how to prevent concurrent updates to shared state
 - Use locks
- At a high-level that is exactly what the database does
 - It *locks* the rows that are part of the transaction
 - After a lock is set, it remains in effect till either the transaction is committed or rolled back
 - No other transaction can access the rows that are locked

Transaction Isolation

- Instead of all-or-nothing isolation, can we do better to improve performance?
- Define different types of locks (isolation levels) that provide different concurrency guarantees
 - Read_uncommitted
 - Read_committed
 - Repeatable_read
 - Serializable
 - Guarantees that there will no dirty reads, no non-repeatable reads, and no phantom reads

Concurrent Transactions: Issues

- Read uncommitted level leads to *Dirty reads*
 - A transaction reads a row from a database table containing uncommitted changes from another transaction.
- Read committed leads to *Non-repeatable reads*
 - A transaction reads a row from a database table, a second transaction changes the same row and the first transaction re-reads the row and gets a different value.
- Repeatable reads leads to *Phantom reads*
 - A transaction re-executes a query, returning a set of rows that satisfies a search condition and finds that another committed transaction has inserted additional rows that satisfy the condition.
- <http://www.onjava.com/pub/a/onjava/2001/05/23/j2ee.html?page=2>

Transaction Isolation Levels

- Read Uncommitted
 - Read data that is not yet committed (dirty reads)
 - Read Committed
 - No dirty reads; but non-repeatable reads
 - Repeatable Read
 - No dirty reads
 - No non-repeatable reads
 - Serializable
 - No dirty reads
 - No non-repeatable reads
 - No phantom reads
- Example:
- `findTransactionIsolationLevel`
 - `testRepeatableRead`
 - `testSerializableRead`

Transaction Isolation Levels

- Serializable
 - <http://sqlperformance.com/2014/04/t-sql-queries/the-serializable-isolation-level>
- Repeatable read
 - <http://sqlperformance.com/2014/04/t-sql-queries/the-repeatable-read-isolation-level>
- Why use read-uncommitted?
 - <http://stackoverflow.com/questions/2471055/why-use-a-read-uncommitted-isolation-level>

JDBC Reference

- <https://github.com/devdattakulkarni/ModernWebApps/tree/master/JDBC>
- <http://stackoverflow.com/questions/4246646/mysql-java-get-id-of-the-last-inserted-value-jdbc>
- <http://stackoverflow.com/questions/42648/best-way-to-get-identity-of-inserted-row>
- <http://stackoverflow.com/questions/8146793/no-suitable-driver-found-for-jdbcmysql-localhost3306-mysql>

JDBC Reference

- <http://stackoverflow.com/questions/9428573/is-it-safe-to-use-a-static-java-sql-connection-instance-in-a-multithreaded-system>
- <http://stackoverflow.com/questions/7592056/am-i-using-jdbc-connection-pooling>
- Apache DBCP2:
 - <http://svn.apache.org/viewvc/commons/proper/dbcp/trunk/doc/BasicDataSourceExample.java?view=markup>

JDBC References

- Auto increment id
 - <http://stackoverflow.com/questions/1915166/how-to-get-the-insert-id-in-jdbc>
- Foreign Key
 - <http://stackoverflow.com/questions/25920251/how-to-automatically-insert-foreign-key-references-in-tables-in-mysql-or-jdbc>
- Triggers
 - <http://dev.mysql.com/doc/refman/5.7/en/trigger-syntax.html>